

Dueling GPUs

Peter Brooker


March 2024

Outline

- Dueling GPU code description
- Sample Game Output
- Kernel Used by Code
- Implementation Details
- Potential Code improvements

Dueling GPU description

Dueling GPUs description



- Assignment was for Dueling GPUs
- Unfortunately, Alienware Minitower has only 1 CUDA GPU
- Code was created to “Virtually” access 2nd GPU
 - Memory Allocation explicitly for d0_... and d1_...
 - If 2nd GPU present, code uses cudaSetDevice(1)
- Uses concept of GPU weighted board vector multiplied by each 1 of 10 solution vectors
- Dot product SUM used to identify 4 in a row as well as next guess
- If GPU#1 is first, “X” placed in Top Left
- Else, GPU#2 “O” placed randomly

```
if (nDevices > 1)
{
    cudaSetDevice(1);
}
```

Sample Game Output

```
*****
* 2 GPUs compete in 4x4 Tic-Tack-Toe *
*****
```

Number of CUDA capable devices = 1

2nd GPU starts

```
- | - | - | 0
---|---|---|---
- | - | - | -
---|---|---|--- End of First Turn
- | - | - | -
---|---|---|---
- | - | - | -
```

Now 1st GPU

```
h_boardOfNums:      0      0      0      1      0      0      0      0      0      0      0      0      0      0      0      0
Vector for SV multiply : 1      1      1    -100    1      1      1      1      1      1      1      1      1      1      1      1
```

```
- | - | - | 0
---|---|---|---
- | - | - | X
---|---|---|--- End of Turn
- | - | - | -
---|---|---|---
- | - | - | -
```

Now 2nd GPU.

h_boardOfNums:	0	0	0	1	0	0	0	-1	0	0	0	0	0	0	0	0
Vector for SV multiply :	1	1	1	10	1	1	1	-100	1	1	1	1	1	1	1	1

-		-		0		0	
---		---		---		---	
-		-		-		X	
---		---		---		---	End of Turn
-		-		-		-	
---		---		---		---	
-		-		-		-	

Now 1st GPU

h_boardOfNums:	0	0	1	1	0	0	0	-1	0	0	0	0	0	0	0	0
Vector for SV multiply :	1	1	-100	-100	1	1	1	10	1	1	1	1	1	1	1	1

-		-		0		0	
---		---		---		---	
-		-		X		X	
---		---		---		---	End of Turn
-		-		-		-	
---		---		---		---	
-		-		-		-	

Now 2nd GPU.

h_boardOfNums:	0	0	1	1	0	0	-1	-1	0	0	0	0	0	0	0	0
Vector for SV multiply :	1	1	10	10	1	1	-100	-100	1	1	1	1	1	1	1	1

-		0		0		0
---		---		---		---
-		-		X		X
---		---		---		---
-		-		-		-
---		---		---		---
-		-		-		-

End of Turn

Now 1st GPU

h_boardOfNums:	0	1	1	1	0	0	-1	-1	0	0	0	0	0	0	0	0
Vector for SV multiply :	1	-100	-100	-100	1	1	10	10	1	1	1	1	1	1	1	1

-		0		0		0
---		---		---		---
-		X		X		X
---		---		---		---
-		-		-		-
---		---		---		---
-		-		-		-

End of Turn

Now 2nd GPU.

h_boardOfNums:	0	1	1	1	0	-1	-1	-1	0	0	0	0	0	0	0	0
Vector for SV multiply :	1	10	10	10	1	-100	-100	-100	1	1	1	1	1	1	1	1

0		0		0		0
---		---		---		---
-		X		X		X
---		---		---		---
-		-		-		-
---		---		---		---
-		-		-		-

Four in a row for 2nd GPU

4 in a row for 2nd GPU
*****end of game*****

D:\Work\CUDA\CudaGPUBattle\x64\Debug\CudaGPUBattle.exe (process 4768) exited with code 0.

Press any key to close this window . . .

Kernel Used by Code:

Cuda Device Kernel for SV array multiplication

```
// same kernel for whether it is on GPU#1 or GPU#2
// the kernel call will pass in different array pointers
// also, the cudaMemcpy() will use different source arrays
__global__ void SVKernel(int* d_forDotProduct, int* d_SV01, int* d_SV02, int* d_SV03, int* d_SV04, int* d_SV05,
int* d_SV06, int* d_SV07, int* d_SV08, int* d_SV09, int* d_SV10,
int* d_dot01, int* d_dot02, int* d_dot03, int* d_dot04, int* d_dot05,
int* d_dot06, int* d_dot07, int* d_dot08, int* d_dot09, int* d_dot10 )
{
    int i = threadIdx.x;
    d_dot01[i] = d_forDotProduct[i] * d_SV01[i];
    d_dot02[i] = d_forDotProduct[i] * d_SV02[i];
    d_dot03[i] = d_forDotProduct[i] * d_SV03[i];
    d_dot04[i] = d_forDotProduct[i] * d_SV04[i];
    d_dot05[i] = d_forDotProduct[i] * d_SV05[i];
    d_dot06[i] = d_forDotProduct[i] * d_SV06[i];
    d_dot07[i] = d_forDotProduct[i] * d_SV07[i];
    d_dot08[i] = d_forDotProduct[i] * d_SV08[i];
    d_dot09[i] = d_forDotProduct[i] * d_SV09[i];
    d_dot10[i] = d_forDotProduct[i] * d_SV10[i];
}
```

- Just performs multiplication with each one of 10 solution vectors
- Sum of multiplication is the vector “dot” product.

Implementation Details:

Solution Vectors (SV) defined in code

- Represents board as single 16 element, scalar array.
- A 4 x 4 board has 10 possible solutions
 - 4 horizontal rows
 - 4 vertical columns
 - 2 diagonals
- Each solution can be represented as 16x1 array of 0s & 1s
 - Top row: {1,1,1,1, 0,0,0,0, 0,0,0,0, 0,0,0,0}
 - 1st column: {1,0,0,0, 1,0,0,0, 1,0,0,0, 1,0,0,0}
 - Diagonal: {1,0,0,0, 0,1,0,0, 0,0,1,0, 0,0,0,1}

10 Solution Vectors (SV) defined in code

- Will multiply SVs by board weighted for each GPU
- Dot product sum will be used for next guess
- Sum also tests if board is at solution

```
// set up solution vectors for 10 possible solution. hz row, vt col, diagonal.  
int h_SV01[16] = { 1,1,1,1, 0,0,0,0, 0,0,0,0, 0,0,0,0 }; // top row  
int h_SV02[16] = { 0,0,0,0, 1,1,1,1, 0,0,0,0, 0,0,0,0 }; // 2nd row  
int h_SV03[16] = { 0,0,0,0, 0,0,0,0, 1,1,1,1, 0,0,0,0 }; // 3rd row  
int h_SV04[16] = { 0,0,0,0, 0,0,0,0, 0,0,0,0, 1,1,1,1 }; // bot row  
int h_SV05[16] = { 1,0,0,0, 1,0,0,0, 1,0,0,0, 1,0,0,0 }; // 1st col  
int h_SV06[16] = { 0,1,0,0, 0,1,0,0, 0,1,0,0, 0,1,0,0 }; // 2nd col  
int h_SV07[16] = { 0,0,1,0, 0,0,1,0, 0,0,1,0, 0,0,1,0 }; // 3rd col  
int h_SV08[16] = { 0,0,0,1, 0,0,0,1, 0,0,0,1, 0,0,0,1 }; // 4th col  
int h_SV09[16] = { 1,0,0,0, 0,1,0,0, 0,0,1,0, 0,0,0,1 }; // NW-SE diag  
int h_SV10[16] = { 0,0,0,1, 0,0,1,0, 0,1,0,0, 1,0,0,0 }; // SW-NE diag
```

..h_boardOfNums array

- GPU#1 “X”: value = -1 in array
- GPU#2 “O”: value = 1 in array
- Open Tile: value = 0 in array

```
- | - | - | 0
- | - | - | X
- | - | - |
- | - | - |
- | - | - |

End of Turn

Now 2nd GPU.
h_boardOfNums:      0      0      0      1      0      0      0      -1      0      0      0      0      0      0      0      0
Vector for SV multiply : 1      1      1     10      1      1      1    -100      1      1      1      1      1      1      1      1
```

GPU weighted board vector

```
- | - | - | O
---|---|---|---
- | - | - | X
---|---|---|--- End of Turn
- | - | - | -
---|---|---|---
- | - | - | -

Now 2nd GPU.
h_boardOfNums:      0      0      0      1      0      0      0      -1      0      0      0      0      0      0      0      0
Vector for SV multiply : 1      1      1     10      1      1      1    -100      1      1      1      1      1      1      1      1
```

- Above is for GPU#2 (“O”) which has value 1 in h_boardOfNums
- GPU#2 Weighted vector =10 if tile is “O”, -100 if “X”, 1 if open

Code for GPU#2 weighted vector

```
cout << "Vector for SV multiply : ";
int h_forDotProduct[16];
for (int j = 0; j < 16; ++j)
{
    if (h_boardOfNums[j] == -1)    //this is GPU0...1st GPU
    {
        h_forDotProduct[j] = -100;
    }
    else if (h_boardOfNums[j] == 1) // this is GPU1...2nd GPU
    {
        h_forDotProduct[j] = 10;
    }
    else if (h_boardOfNums[j] == 0)
    {
        h_forDotProduct[j] = 1;
    }
}
```

Now 2nd GPU.

h_boardOfNums:	0	0	0	1	0	0	0	-1	0	0	0	0	0	0	0	0
Vector for SV multiply :	1	1	1	10	1	1	1	-100	1	1	1	1	1	1	1	1

Next Guess Identification

- On host side, create sums of all d_dotXX arrays returned from Kernel
 - $XX = 01, 02, \dots, 10$
 - Sum is the “dot product” of GPU weighted board array with each of the solution vectors
- Identify the max sum
- If max sum = $10 + 10 + 10 + 1 = 31$, next guess is in the open space and game is over

Potential Code Improvements

Potential Code Improvements

- Guessing
- Proposed Improvement:
 - Place all code in functions after main
 - In main, only highlevel functions are called
 - Function names would be very descriptive so main would read like a flow chart
 - Why not done? Returning arrays defined in functions confusing due to heap vs stack persistence issues.

Thank You!